

Rq2 reproduction

Jeffrey M. Young

15 May 2020

In this walk-through we'll perform the plot generation and significance testing for rq2. We assume you have read the `rq1_rq3` and `sigTest` walk-throughs, in that order.

Preliminaries

First, libraries, note that I have silenced the shadowing warnings from R:

```
library(ggplot2)    #plotting
library(dplyr)      #dataframe manipulation
library(tidyr)      #dataframe manipulation
library(broom)      #for the tidy function
library(scales)     #for scientific function
library(latex2exp)  #for TeX in Labels
library(ggpubr)     #for stat_cor
```

Load the data, and perform some simple munging to select relevant columns

```
finResultsFile <- "../../../data/fin_comp_data.csv"
autoResultsFile <- "../../../data/auto_comp_data.csv"

finData <- read.csv(file=finResultsFile) %>%
  mutate(Algorithm = as.factor(Algorithm), Config = as.factor(Config))

autoData <- read.csv(file=autoResultsFile) %>%
  mutate(Algorithm = as.factor(Algorithm), Config = as.factor(Config))

finDF <- finData %>% mutate(data = "Fin") %>%
  select(Mean, Algorithm, CompressionRatio, data, ChcCount, PlainCount, Config)

autoDF <- autoData %>% mutate(data = "Auto") %>%
  select(Mean, Algorithm, CompressionRatio, data, ChcCount, PlainCount, Config)
```

Visualization

Now the sharing ratio must be calculated, this is performed via a `spread` and `mutate` operation. `spread` transforms long data into wide data, note that it has been recently deprecated and replaced with `pivot_wider`, see the `tidyverse` documentation for details here. We use `spread` just for convenience and familiarity. Now we find the sharing ratio for each Algorithm, and construct a data frame for each dataset:

```
## finding the sharing ratios for auto
autoPtoP <- autoDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "p-->p") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `p-->p`)) %>%
  select(-"v-->p", -"p-->p") %>% mutate(Algorithm = "p\u27f6p")
```

```

autoVtoV <- autoDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "v-->v") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `v-->v`)) %>%
  select("-v-->v", "-v-->p") %>% mutate(Algorithm = "v\u27f6v")

autoPtoV <- autoDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "p-->v") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `p-->v`)) %>%
  select("-v-->p", "-p-->v") %>% mutate(Algorithm = "p\u27f6v")

## construct auto data frame
autoSharedDF <- rbind(autoPtoP, autoVtoV, autoPtoV)

```

and for financial:

```

## finding the sharing ratios for fin
finPtoP <- finDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "p-->p") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `p-->p`)) %>%
  select("-v-->p", "-p-->p") %>% mutate(Algorithm = "p\u27f6p")

finVtoV <- finDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "v-->v") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `v-->v`)) %>%
  select("-v-->v", "-v-->p") %>% mutate(Algorithm = "v\u27f6v")

finPtoV <- finDF %>% group_by(Algorithm, Config) %>%
  filter(Algorithm == "v-->p" || Algorithm == "p-->v") %>%
  spread(Algorithm, Mean) %>% mutate(MeanRatio = (`v-->p` / `p-->v`)) %>%
  select("-v-->p", "-p-->v") %>% mutate(Algorithm = "p\u27f6v")

## construct fin data frame
finSharedDF <- rbind(finPtoP, finVtoV, finPtoV)

```

Now the data frames are combined and the plot can be generated:

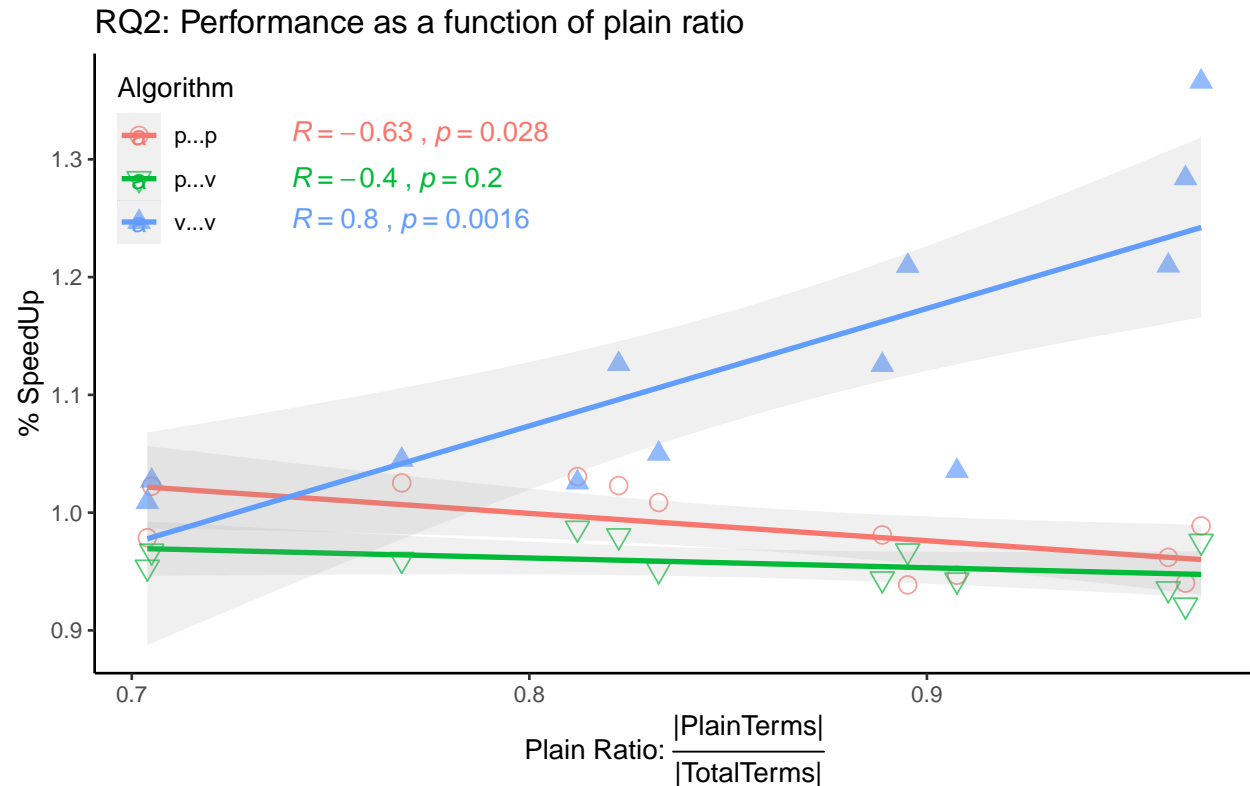
```

## make the data frame
df <- rbind(autoSharedDF, finSharedDF) %>%
  ## calculate the ratio of plain terms to total terms, i.e., the sharing ratio
  mutate(PlainRatio = PlainCount / (ChcCount + PlainCount),
  ## get the converse ratio at the same time
  ChcRatio = ChcCount / (ChcCount + PlainCount))

## make the plot as a function of plain ratio to mean ratio
ggplot(df, mapping = aes(x=PlainRatio, y=MeanRatio, colour = Algorithm, shape = Algorithm)) +
  ## it is a scatter plot
  geom_point(size=3, alpha=0.7) +
  ylab("% SpeedUp") +
  ## insert LaTeX into the x-axis labels
  xlab(TeX("Plain Ratio:  $\frac{|Plain Terms|}{|Total Terms|}$ ")) +
  ## manually scale shapes so the algorithms have the same shape over different plots
  scale_shape_manual(values = c(1,6,17)) +
  ## add the confidence intervals, using a linear regression, alpha makes these more
  ## transparent, se adds the shadings for the 95% confidence intervals
  geom_smooth(method=lm, formula = y ~ x, se=TRUE, alpha=0.15) +

```

```
ggtitle("RQ2: Performance as a function of plain ratio") +
theme_classic() +
## stat_cor calculates the R^2 values, that are placed next to the legend
stat_cor(aes(color=Algorithm),
         label.x=0.74, label.y.npc=c(0.91, 0.89, 0.88)) +
theme(legend.position = c(0.07, 0.83))
```



Statistical Significance

We perform the same two way ANOVA to ensure fair comparisons:

```
### Perform the anova
res.aov <- aov(MeanRatio ~ PlainRatio * Algorithm, data = df)
res.aov

## Call:
## aov(formula = MeanRatio ~ PlainRatio * Algorithm, data = df)
##
## Terms:
##          PlainRatio  Algorithm PlainRatio:Algorithm  Residuals
## Sum of Squares  0.01522695  0.19480856          0.08822280  0.06465027
## Deg. of Freedom           1           2              2          30
##
## Residual standard error: 0.04642207
## Estimated effects may be unbalanced

### Check the summary to see what is significant, all of it is as expected
summary(res.aov)

##              Df Sum Sq Mean Sq F value    Pr(>F)
```

```
## PlainRatio          1 0.01523 0.01523  7.066  0.0125 *
## Algorithm           2 0.19481 0.09740 45.199 8.86e-10 ***
## PlainRatio:Algorithm 2 0.08822 0.04411 20.469 2.48e-06 ***
## Residuals          30 0.06465 0.00216
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

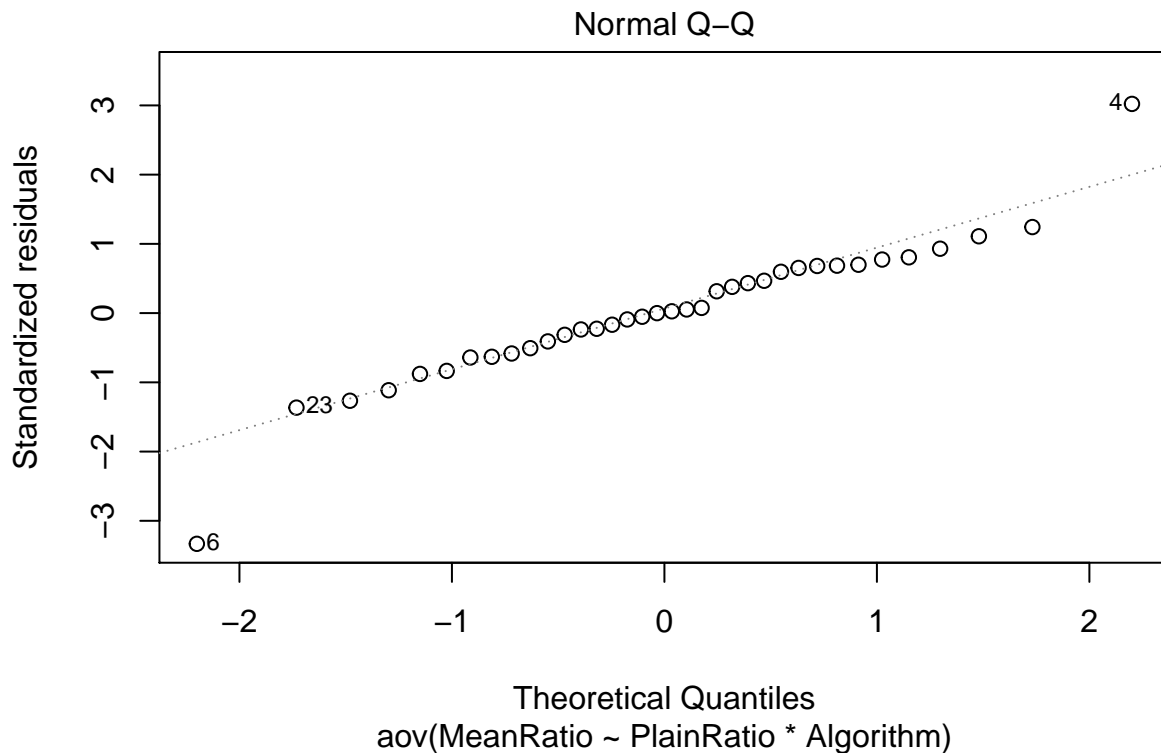
We see that Algorithm, PlainRatio, and their interaction are significant. This makes sense given the nature of the tool and analysis. We perform a Tukey pair-wise comparison to check the significantly different comparisons in the dataset.

```
### perform the pair-wise Tukey comparison to test the difference
### between groups
TukeyHSD(res.aov) %>% tidy %>% mutate(pVal = scientific(adj.p.value, 3))
```

```
## # A tibble: 3 x 7
##   term      comparison estimate conf.low conf.high  adj.p.value pVal
##   <chr>      <chr>      <dbl>   <dbl>   <dbl>      <dbl> <chr>
## 1 Algorithm p v-p p    -0.0300 -0.0768  0.0167 0.268      2.68e-01
## 2 Algorithm v v-p p     0.139    0.0921  0.186 0.000000109 1.09e-07
## 3 Algorithm v v-p v     0.169    0.122   0.216 0.00000000185 1.85e-09
```

We observe that v-->v differs from the other algorithms, and the difference is statistically significant. Next, we check the residuals to assess normality.

```
res.ass <- plot(res.aov, 2)
```



Residuals look good and there are three outliers, which agrees with other analyses. To verify, we perform the Shapiro-Wilks test, identically to `rq3`.

```
aov.resids <- residuals(object=res.aov)
```

```
## Shapiro-Wilk normality test
```

```
shapiro.test(x = aov.resids)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  aov.resids  
## W = 0.92064, p-value = 0.01315
```

We see that again the dataset violates the assumptions of the ANOVA test, so we perform the Kruskal-Wallis test instead:

```
## Algorithms are significant  
kruskal.test(MeanRatio ~ Algorithm, df)
```

```
##  
## Kruskal-Wallis rank sum test  
##  
## data:  MeanRatio by Algorithm  
## Kruskal-Wallis chi-squared = 23.73, df = 2, p-value = 7.033e-06
```

Algorithms are found to be statistically different as expected.

```
## Plain ratio is not significant!!  
kruskal.test(MeanRatio ~ PlainRatio, df)
```

```
##  
## Kruskal-Wallis rank sum test  
##  
## data:  MeanRatio by PlainRatio  
## Kruskal-Wallis chi-squared = 4.2733, df = 11, p-value = 0.9612
```

Plain ratio, interestingly is found not to be statistically different.

```
## Interaction is not significant surprisingly  
rq2.inters <- interaction(df$Algorithm, df$PlainRatio)  
kruskal.test(MeanRatio ~ rq2.inters, df)
```

```
##  
## Kruskal-Wallis rank sum test  
##  
## data:  MeanRatio by rq2.inters  
## Kruskal-Wallis chi-squared = 35, df = 35, p-value = 0.4682
```

However, when accounting for Algorithms, i.e., when analyzing the interaction between the two, we find the interaction is significant

```
## Show the pairs which are significant  
pairs <- pairwise.wilcox.test(df$MeanRatio, df$Algorithm,  
                             p.adj="bonf" , exact=TRUE, paired=FALSE) %>%  
  tidy %>% arrange(p.value)
```

```
pairs
```

```
## # A tibble: 3 x 3  
##   group1 group2   p.value  
##   <chr>  <chr>     <dbl>  
## 1 v v    p v      0.00000222  
## 2 v v    p p       0.0000666  
## 3 p v    p p       0.116
```